

Lab1文档

“任务与思考题”部分仅直截了当记录实验结果并回答问题，具体实验细节在“实验过程”部分

任务与思考题

1. 内存分配大小分析

请求分配大小	实际分配大小	分析与说明
96 字节	128字节	96+32 (meta大小) <819字节，使用salloc进行分配，向上取整分配1*128字节的小页面
128 字节	256字节	128+32 (meta大小) <819字节，使用salloc进行分配，向上取整分配2*128字节的小页面
256 字节	384字节	256+32 (meta大小) <819字节，使用salloc进行分配，向上取整分配3*128字节的小页面
4064 字节	4096字节	4064+32 (meta大小) >819字节，使用palloc进行分配，向上取整分配1*4096字节的页面
4096 字节	8192字节	4096+32 (meta大小) >819字节，使用palloc进行分配，向上取整分配2*4096字节的页面

2. 核心问题

(1) 最小分配单元: Unikraft 两种内存分配策略的最小单元是多少？它是如何定义的？

Unikraft采用salloc和palloc两种内存分配策略，其中：

- salloc分配的最小单元为128字节，在__S_PAGE_SIZE常量中定义，salloc时根据请求大小向上取整分配若干128字节大小的页面
- palloc分配的最小单元为4096字节，在__PAGE_SIZE常量中定义，palloc时根据请求大小向上取整分配若干4096字节大小的页面

(2) 分配器选择：uk_malloc() 函数在何种条件下会选择 palloc，又在何种条件下会选择salloc？

在处理大小为size的分配请求时，首先计算size+METADATA_IFPAGES_SIZE_POW2（meta数据的大小）是否<__PAGE_SIZE/5（819），如果分配请求真实需要分配的大小<819个字节则使用salloc进行分配，否则使用palloc进行分配

(3) 大内存分配问题: 当前 palloc 在处理大内存（例如，一次性分配多个页面）的分配与回收时，存在一个已知的设计问题。请定位该问题，并尝试在 GDB 中通过 set 命令修改相关变量，模拟正确的 free 过程，并截图记录结果。

问题出现在small参数没有被正确进行初始化（是一个很奇怪的指针值），导致在根据small参数判断“究竟该使用大页面的释放方式还是小页面的释放方式”时，模型总是会选择按照小页面的方式进行释放，导致使用palloc申请的内存区域无法正确释放。共做出两处改动：

1.在alloc.c设置is_small参数为0后，成功调用大页面的PAGE_ALIGN_DOWN函数而不是小页面的PAGE_ALIGN_DOWN_S

```
-- at /workspaces/OS-2026/app-helloworld/workdir/unikraft/lib/ukalloc/alloc.c:116
116     UK_ASSERT((__uptr)ptr >= __PAGE_SIZE + METADATA_IFPAGES_SIZE_POW2);
(gdb) s
119     int is_small = (int)small;
(gdb) s
120     if (is_small) {
(gdb) set variable is_small = 0
(gdb) p is_small
$1 = 0
(gdb) s
130     metadata = PAGE_ALIGN_DOWN((__uptr)ptr);
(gdb) s
131     if (metadata == (__uptr)ptr) {
```

2.在判断sfree和pfree之前手动设置small参数为0，成功调用uk_pfree函数而不是uk_sfree函数进行内存释放

```
(gdb) p *small
Attempt to dereference a generic pointer.
(gdb) set variable small = 0
(gdb) p small
$4 = (const void *) 0x0
(gdb) s
224     uk_pfree(a, metadata->base, metadata->num_pages);
(gdb) s
uk_pfree (a=0x40010000, ptr=0x40192000, num_pages=1)
    at /workspaces/OS-2026/app-helloworld/workdir/unikraft/lib/ukalloc/include/uk/alloc.h:249
249     uk_do_pfree(a, ptr, num_pages);
(gdb) s
uk_do_pfree (a=0x40010000, ptr=0x40192000, num_pages=1)
    at /workspaces/OS-2026/app-helloworld/workdir/unikraft/lib/ukalloc/include/uk/alloc.h:238
238     UK_ASSERT(a);
```

关于手动设置is_small参数之后，free函数的表现差异（内存变化结果）呈现在实验过程章节

实验过程

1. 内存分配大小分析

(1) 在main.c malloc和free的位置设置断点，对于malloc其调用链可追溯为：malloc -> uk_malloc -> uk_do_malloc -> uk_malloc_ifpages；对于uk_malloc_ifpages函数，对其注释分析如下：

代码块

```
1 void *uk_malloc_ifpages(struct uk_alloc *a, __sz size) {
2     __uptr intptr;
3     unsigned long num_pages;
4     struct metadata_ifpages *metadata;
5     #ifdef CONFIG_HAVE_MEMTAG
6         size = MEMTAG_ALIGN(size);
7     #endif
8     /*由于需要空间存储分配内存块的元数据信息，因而这里需要添加一个固定大小的常量
9      METADATA_IFPAGES_SIZE_POW2 (值为32) */
10    __sz realsize = METADATA_IFPAGES_SIZE_POW2 + size;
11    UK_ASSERT(a);
12    if (!size || realsize < size) return __NULL;
13    /*这里使用IS_SMALL函数判断真实待分配区域的大小是否<_PAGE_SIZE/5，其中_PAGE_SIZE为
14     4096，所以临界点为819字节，大于等于这个临界点使用palloc，否则使用salloc*/
15    if (IS_SMALL(realsize)) {
16        /*判断需要几个128字节的小页面，并调用uk_salloc函数进行分配*/
17        num_pages = size_to_s_num_pages(realsize);
18        intptr = (__uptr)uk_salloc(a, num_pages);
19        uk_pr_err("alloc size => %llu, num_pages => %llu, intptr => %p\n",
20                  realsize, num_pages, intptr);
21    }
22    /*判断需要几个4096字节的大页面，并调用uk_palloc函数进行分配*/
23    else {
24        num_pages = size_to_num_pages(realsize);
25        intptr = (__uptr)uk_palloc(a, num_pages);
26        if (!intptr) return __NULL;
27        /*这里在最开始32个字节内的区域设置元信息，包括大小、起始地址和分配的页数*/
28        metadata = (struct metadata_ifpages *)intptr;
29        metadata->size = size;
30        metadata->num_pages = num_pages;
31        metadata->base = (void *)intptr;
32    #ifdef CONFIG_HAVE_MEMTAG
33        return ukarch_memtag_region((void *)(intptr + METADATA_IFPAGES_SIZE_POW2),
34                                     size);
35    #else
36        /*由于分配区域前32个字节是meta数据，因而程序实际可用的是32字节以后的位置*/
37        return (void *)(intptr + METADATA_IFPAGES_SIZE_POW2);
38    #endif
39 }
```

```
39 }
```

(2) 在对调用malloc/palloc进行内存分配后，我们可以使用gdb查看各内存块metadata的情况：

对于malloc的情况（96/128/256字节）：

```
(gdb) p *metadata
$6 = {size = 96, num_pages = 1, base = 0x44101a00}

199      return (void *) (intptr + METADATA_IFPAGES_SIZE_POW2);
(gdb) p *metadata
$3 = {size = 128, num_pages = 2, base = 0x44101980}

(gdb) p *metadata
$2 = {size = 256, num_pages = 3, base = 0x44101900}
```

对于palloc的情况（4064/4096字节）：

```
(gdb) p *metadata
$2 = {size = 4064, num_pages = 1, base = 0x40192000}

199      return (void *) (intptr + METADATA_IFPAGES_SIZE_POW2);
(gdb) p *metadata
$2 = {size = 4096, num_pages = 2, base = 0x40192000}
```

(3) 同样地，对在free位置的断点进行深入，直到uk_free_ifpages函数，对该函数进行注释分析如下：

代码块

```
1 void uk_free_ifpages(struct uk_alloc *a, void *ptr, const void *small) {
2     struct metadata_ifpages *metadata;
3 #ifdef CONFIG_HAVE_MEMTAG
4     __sz size;
5 #endif
6     UK_ASSERT(a);
7     if (!ptr) return;
8 #ifdef CONFIG_HAVE_MEMTAG
9     metadata = uk_get_metadata((void *)((__u64)ptr & ~MTE_TAG_MASK));
10    size = metadata->size;
11 #else
12     /*获得该内存块的元数据信息（页数、起始地址等）*/
13     metadata = uk_get_metadata(ptr, small);
14 #endif
15     UK_ASSERT(metadata->base != __NULL);
16     UK_ASSERT(metadata->num_pages != 0);
17     /*如果是128字节的小页面，调用uk_sfree释放metadata中记录的num_pages个页面*/
18     if (small) {
19         uk_sfree(a, metadata->base, metadata->num_pages);
20     }
21     /*否则调用uk_pfree释放记录的num_pages个4096字节大小的正常页面*/
```

```

22     else {
23         uk_pfree(a, metadata->base, metadata->num_pages);
24     }
25
26 #ifdef CONFIG_HAVE_MEMTAG
27     ukarch_memtag_region(ptr, size);
28 #endif
29 }
```

2.free算法流程模拟

理论上，在调用malloc(4096)前和free之后，内存空闲情况应保持一致

在未改动代码情况下，由于small参数未正确初始化，导致内存申请释放前后系统内存状态不一致（左边为申请前，右边为释放后）

代码块

```

1 [ 4.239707] ERR:
[libukallocbbuddy] <bbuddy.c @
305> Dumping current state of
the free list:
2 [ 4.240418] ERR:
[libukallocbbuddy] <bbuddy.c @
320> Free list 0 is empty.
3 [ 4.240992] ERR:
[libukallocbbuddy] <bbuddy.c @
322> Free list 1 (Order 1):
4 [ 4.241536] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x40192000, end: 0x40194000,
level: 1
5 [ 4.242440] ERR:
[libukallocbbuddy] <bbuddy.c @
322> Free list 2 (Order 2):
6 [ 4.242999] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x44104000, end: 0x44108000,
level: 2
7 [ 4.243844] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
```

代码块

```

1 [ 4.292088] ERR:
[libukallocbbuddy] <bbuddy.c @
345> Dumping current state of
the sfree list:
2 [ 4.292664] ERR:
[libukallocbbuddy] <bbuddy.c @
350> SFree list 0 (Order 0):
3 [ 4.293150] ERR:
[libukallocbbuddy] <bbuddy.c @
354> Entry at address:
0x44101a00, end: 0x44101a80,
level: 0
4 [ 4.293851] ERR:
[libukallocbbuddy] <bbuddy.c @
354> Entry at address:
0x44101980, end: 0x44101a00,
level: 0
5 [ 4.294493] ERR:
[libukallocbbuddy] <bbuddy.c @
354> Entry at address:
0x44101900, end: 0x44101980,
level: 0
6 [ 4.294906] ERR:
[libukallocbbuddy] <bbuddy.c @
354> Entry at address:
0x44101880, end: 0x44101900,
level: 0
```

```
0x40194000, end: 0x40198000,  
level: 2  
8 [ 4.244281] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40014000, end: 0x40018000,  
level: 2  
9 [ 4.244733] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 3 (Order 3):  
10 [ 4.245043] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44108000, end: 0x44110000,  
level: 3  
11 [ 4.245496] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40198000, end: 0x401a0000,  
level: 3  
12 [ 4.245893] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40018000, end: 0x40020000,  
level: 3  
13 [ 4.246510] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 4 (Order 4):  
14 [ 4.247293] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44110000, end: 0x44120000,  
level: 4  
15 [ 4.248391] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 5 is empty.  
16 [ 4.249143] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 6 (Order 6):  
17 [ 4.249873] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44140000, end: 0x44180000,  
level: 6  
18 [ 4.250499] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:
```

```
7 [ 4.295382] ERR:  
[libukallocbbuddy] <bbuddy.c @  
354> Entry at address:  
0x44101800, end: 0x44101880,  
level: 0  
8 [ 4.295789] ERR:  
[libukallocbbuddy] <bbuddy.c @  
354> Entry at address:  
0x44101780, end: 0x44101800,  
level: 0  
9 [ 4.296604] ERR:  
[libukallocbbuddy] <bbuddy.c @  
354> Entry at address:  
0x44101700, end: 0x44101780,  
level: 0  
10 [ 4.297248] ERR:  
[libukallocbbuddy] <bbuddy.c @  
354> Entry at address:  
0x44101680, end: 0x44101700,  
level: 0  
11 [ 4.298089] ERR:  
[libukallocbbuddy] <bbuddy.c @  
354> Entry at address:  
0x44101600, end: 0x44101680,  
level: 0  
12 [ 4.298583] ERR:  
[libukallocbbuddy] <bbuddy.c @  
354> Entry at address:  
0x44101580, end: 0x44101600,  
level: 0
```

```

0x401c0000, end: 0x40200000,
level: 6
19 [ 4.251346] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x40040000, end: 0x40080000,
level: 6
20 [ 4.253035] ERR:
[libukallocbbuddy] <bbuddy.c @
322> Free list 7 (Order 7):
21 [ 4.253628] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x44180000, end: 0x44200000,
level: 7
22 [ 4.254288] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x40080000, end: 0x40100000,
level: 7
23 [ 4.256023] ERR:
[libukallocbbuddy] <bbuddy.c @
320> Free list 8 is empty.
24 [ 4.256357] ERR:
[libukallocbbuddy] <bbuddy.c @
322> Free list 9 (Order 9):
25 [ 4.256685] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x44200000, end: 0x44400000,
level: 9
26 [ 4.257147] ERR:
[libukallocbbuddy] <bbuddy.c @
326> Entry at address:
0x40200000, end: 0x40400000,
level: 9

```

在使用set命令模拟正确行为（修正small参数之后），内存申请释放前后系统内存状态一致（左边为申请前，右边为申请后）

代码块

```

1 [ 4.276722] ERR:
[libukallocbbuddy] <bbuddy.c @
305> Dumping current state of
the free list:

```

代码块

```

1 [ 4.344764] ERR:
[libukallocbbuddy] <bbuddy.c @
305> Dumping current state of
the free list:

```

```
2 [ 4.277390] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 0 is empty.  
3 [ 4.277991] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 1 (Order 1):  
4 [ 4.278576] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40192000, end: 0x40194000,  
level: 1  
5 [ 4.279332] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 2 (Order 2):  
6 [ 4.279942] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44104000, end: 0x44108000,  
level: 2  
7 [ 4.280778] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40194000, end: 0x40198000,  
level: 2  
8 [ 4.281586] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40014000, end: 0x40018000,  
level: 2  
9 [ 4.282453] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 3 (Order 3):  
10 [ 4.283121] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44108000, end: 0x44110000,  
level: 3  
11 [ 4.283618] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40198000, end: 0x401a0000,  
level: 3  
12 [ 4.284389] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40018000, end: 0x40020000,  
level: 3
```

```
2 [ 4.345266] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 0 is empty.  
3 [ 4.345655] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 1 (Order 1):  
4 [ 4.346250] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40192000, end: 0x40194000,  
level: 1  
5 [ 4.346900] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 2 (Order 2):  
6 [ 4.347215] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44104000, end: 0x44108000,  
level: 2  
7 [ 4.348053] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40194000, end: 0x40198000,  
level: 2  
8 [ 4.348734] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40014000, end: 0x40018000,  
level: 2  
9 [ 4.349433] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 3 (Order 3):  
10 [ 4.349767] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44108000, end: 0x44110000,  
level: 3  
11 [ 4.350203] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40198000, end: 0x401a0000,  
level: 3  
12 [ 4.350842] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40018000, end: 0x40020000,  
level: 3
```

```
13 [ 4.285120] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 4 (Order 4):  
14 [ 4.285793] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44110000, end: 0x44120000,  
level: 4  
15 [ 4.286666] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 5 is empty.  
16 [ 4.287190] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 6 (Order 6):  
17 [ 4.287674] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44140000, end: 0x44180000,  
level: 6  
18 [ 4.288478] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x401c0000, end: 0x40200000,  
level: 6  
19 [ 4.289307] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40040000, end: 0x40080000,  
level: 6  
20 [ 4.290069] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 7 (Order 7):  
21 [ 4.290377] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44180000, end: 0x44200000,  
level: 7  
22 [ 4.290825] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40080000, end: 0x40100000,  
level: 7  
23 [ 4.291859] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 8 is empty.  
24 [ 4.292536] ERR:  
[libukallocbbuddy] <bbuddy.c @
```

```
13 [ 4.351791] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 4 (Order 4):  
14 [ 4.352292] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44110000, end: 0x44120000,  
level: 4  
15 [ 4.353012] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 5 is empty.  
16 [ 4.353387] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 6 (Order 6):  
17 [ 4.353915] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44140000, end: 0x44180000,  
level: 6  
18 [ 4.354618] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x401c0000, end: 0x40200000,  
level: 6  
19 [ 4.355062] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40040000, end: 0x40080000,  
level: 6  
20 [ 4.355499] ERR:  
[libukallocbbuddy] <bbuddy.c @  
322> Free list 7 (Order 7):  
21 [ 4.356192] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44180000, end: 0x44200000,  
level: 7  
22 [ 4.357264] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40080000, end: 0x40100000,  
level: 7  
23 [ 4.358048] ERR:  
[libukallocbbuddy] <bbuddy.c @  
320> Free list 8 is empty.  
24 [ 4.358623] ERR:  
[libukallocbbuddy] <bbuddy.c @
```

```
322> Free list 9 (Order 9):  
25 [ 4.293136] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44200000, end: 0x44400000,  
level: 9  
26 [ 4.293952] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40200000, end: 0x40400000,  
level: 9
```

```
322> Free list 9 (Order 9):  
25 [ 4.359185] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x44200000, end: 0x44400000,  
level: 9  
26 [ 4.359941] ERR:  
[libukallocbbuddy] <bbuddy.c @  
326> Entry at address:  
0x40200000, end: 0x40400000,  
level: 9
```